

DATA PARALLELISM

LEARNING OBJECTIVES

- Learn about task parallelism and data parallelism
- Learn about the SPMD model for describing data parallelism
- Learn about SYCL execution and memory models
- Learn about enqueueing kernel functions with `parallel_for`

TASK VS DATA PARALLELISM



- **Task parallelism** is where you have several, possibly distinct tasks executing in parallel.
 - In task parallelism you generally optimize for latency.
- **Data parallelism** is where you have the same task being performed on multiple elements of data.
 - In data parallelism you generally optimize for throughput.

VECTOR PROCESSORS

- Many processors are vector processors, which means they can naturally perform data parallelism.
 - GPUs are designed to be parallel.
 - CPUs have SIMD instructions which perform the same instruction on a number elements of data.

SPMD MODEL FOR DESCRIBING DATA PARALLELISM

Sequential CPU code


```
void calc(const int in[], int out[]) {  
    // all iterations are run in the same  
    // thread in a loop  
    for (int i = 0; i < 1024; i++){  
        out[i] = in[i] * in[i];  
    }  
}  
  
// calc(...) is invoked just once and all  
// iterations are performed inline  
calc(in, out);
```

Parallel SPMD code

```
void calc(const int in[], int out[], int id) {  
    // function is described in terms of  
    // a single iteration  
    out[id] = in[id] * in[id];  
}  
  
// parallel_for invokes calc(...) multiple  
// times in parallel  
parallel_for(calc, in, out, 1024);
```

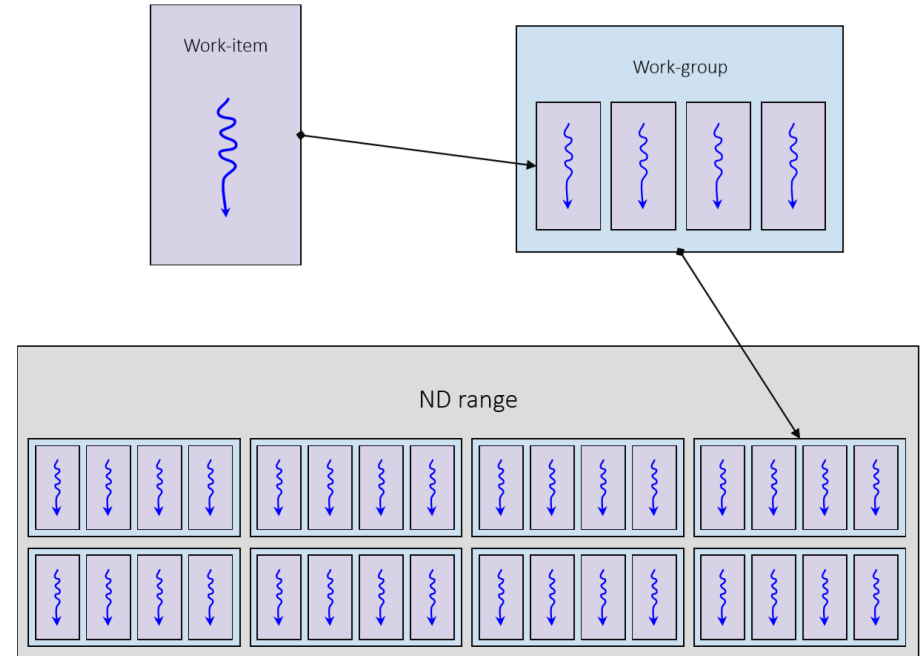
SYCL EXECUTION MODEL

- In SYCL kernel functions are executed by **work-items**.
- You can think of a work-item as a thread of execution.
- Each work-item will execute a SYCL kernel function from start to end.
- A work-item can run on CPU threads, SIMD lanes, GPU threads, or any other kind of processing element.



SYCL EXECUTION MODEL


- SYCL kernel functions are invoked within an **nd-range**
- An nd-range has a number of work-groups and subsequently a number of work-items
- Work-groups always have the same number of work-items



SYCL EXECUTION MODEL

- The nd-range describes an **iteration space**: how it is composed in terms of work-groups and work-items
- An nd-range can have 1, 2 or 3 dimensions
- An nd-range has two components
 - The **global-range** describes the total number of work-items in each dimension
 - The **local-range** describes the number of work-items in a work-group in each dimension


nd-range $\{\{12, 12\}, \{4, 4\}\}$



SYCL EXECUTION MODEL

- Each invocation in the iteration space of an nd-range is a work-item
- Each invocation knows which work-item it is on and can query certain information about its position in the nd-range
- Each work-item has the following:
 - **Global range:** {12, 12}
 - **Global id:** {5, 6}
 - **Group range:** {3, 3}
 - **Group id:** {1, 1}
 - **Local range:** {4, 4}
 - **Local id:** {1, 2}

nd-range {{12, 12}, {4, 4}}




SYCL EXECUTION MODEL

Typically an nd-range invocation SYCL will execute the SYCL kernel function on a very large number of work-items, often in the thousands




SYCL EXECUTION MODEL

- Multiple work-items will generally execute concurrently
- On vector hardware this is often done in lock-step, with neighboring work-items executing the same instruction at all times
- The number of work-items that will execute concurrently can vary from one device to another
- Work-items will be batched along with other work-items in the same work-group
- The order work-items and work-groups are executed in is implementation defined




SYCL EXECUTION MODEL

- Work-items in a work-group can be synchronized using a work-group barrier
 - All work-items within a work-group must reach the barrier before any can continue on




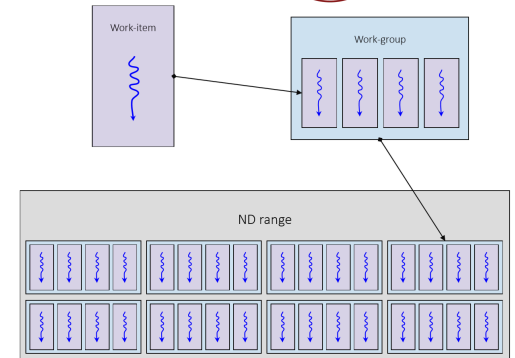
SYCL EXECUTION MODEL

- SYCL does not support synchronizing across all work-items in the nd-range
- The only way to do this is to split the computation into separate SYCL kernel functions



SYCL EXECUTION MODEL

- SYCL also provides a simplified execution model with `sycl::range` in place of `sycl::nd_range`
- Caller only provides the global range
- Local range is decided by the runtime and cannot be inspected
- No synchronization is possible between work items
- Useful for simple problems which don't require synchronization, local memory and ultimate performance
 - Runtime may not always have enough information to choose the best-performing size



PARALLEL_FOR

```
cgh.parallel_for<my_kernel>(sycl::nd_range{{1024, 16}, {32, 4}},  
                             [=](sycl::nd_item<2> item){  
    // SYCL kernel function is executed  
    // on a range of work-items  
});
```

- In SYCL, kernel functions can be enqueued to execute over a range of work-items using `sycl::handler::parallel_for`
- The first argument to `parallel_for` is an `nd_range` or a range which describes the iteration space over which the kernel is to be executed
- The kernel function has to take an `nd_item` or `item`, respectively, as the parameter (or any type they can be implicitly converted to, commonly from `item` to `id`)

EXPRESSING PARALLELISM

```
cgh.parallel_for<kernel>((sycl::nd_range<1>{1024,  
    [=] (sycl::nd_item<1> ndItem) {  
        /* kernel function code */  
        id globalId = ndItem.get_global_id();  
        id localId = ndItem.get_local_id();  
    });
```


```
cgh.parallel_for<kernel>(sycl::range<1>{1024},  
    [=] (sycl::item<1> item) {  
        /* kernel function code */  
        id globalId = item.get_id();  
    });
```

```
cgh.parallel_for<kernel>(sycl::range<1>{1024},  
    [=] (sycl::id<1> globalId) {  
        /* kernel function code */  
    });
```


- Overload taking an `nd_range` object specifies the global and local range
- An `nd_item` parameter represents the global and local range and index
- Overload taking a `range` object specifies the global range, runtime decides local range
- An `item` parameter represents the global range and the index within the global range
- Overload taking a `range` object specifies the global range, runtime decides local range
- An `id` parameter represents the index within the global range

SYCL MEMORY MODEL

- Each work-item can access a dedicated region of **private memory**
- A work-item cannot access the private memory of another work-item

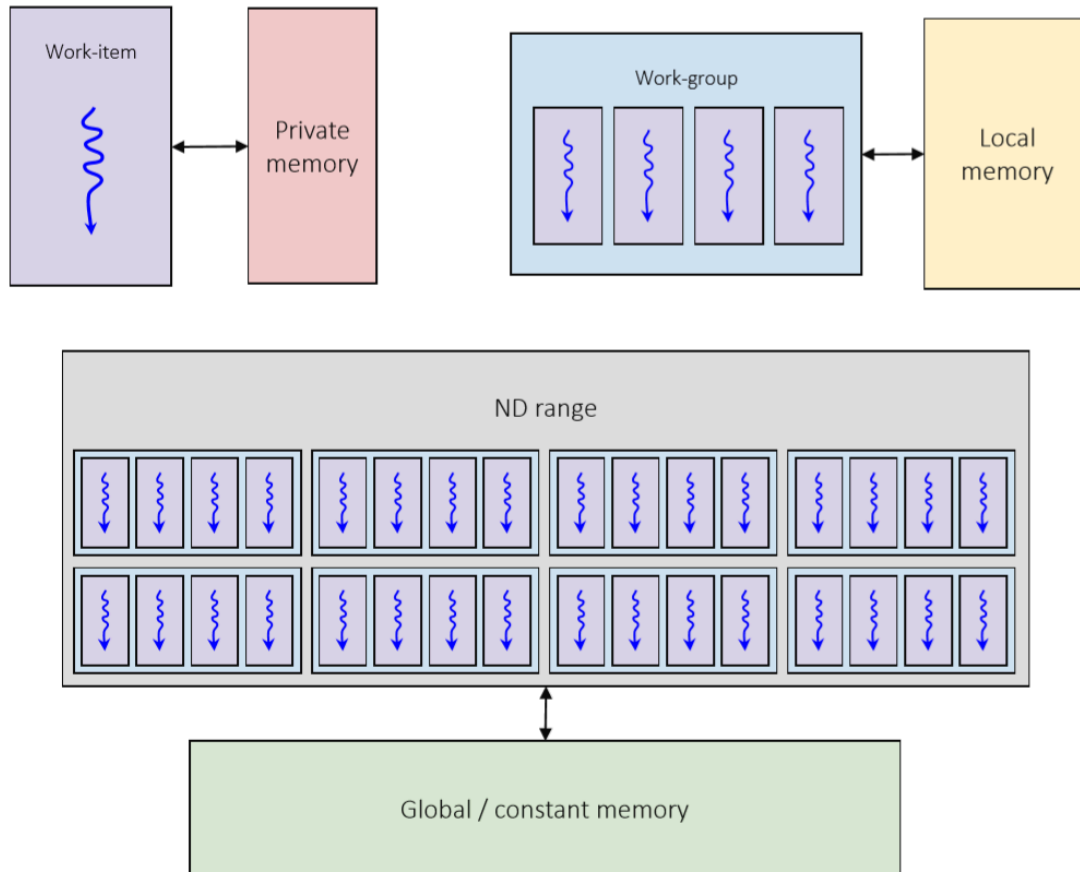


SYCL MEMORY MODEL



- Each work-item can access a dedicated region of **local memory** accessible to all work-items in a work-group
- A work-item cannot access the local memory of another work-group

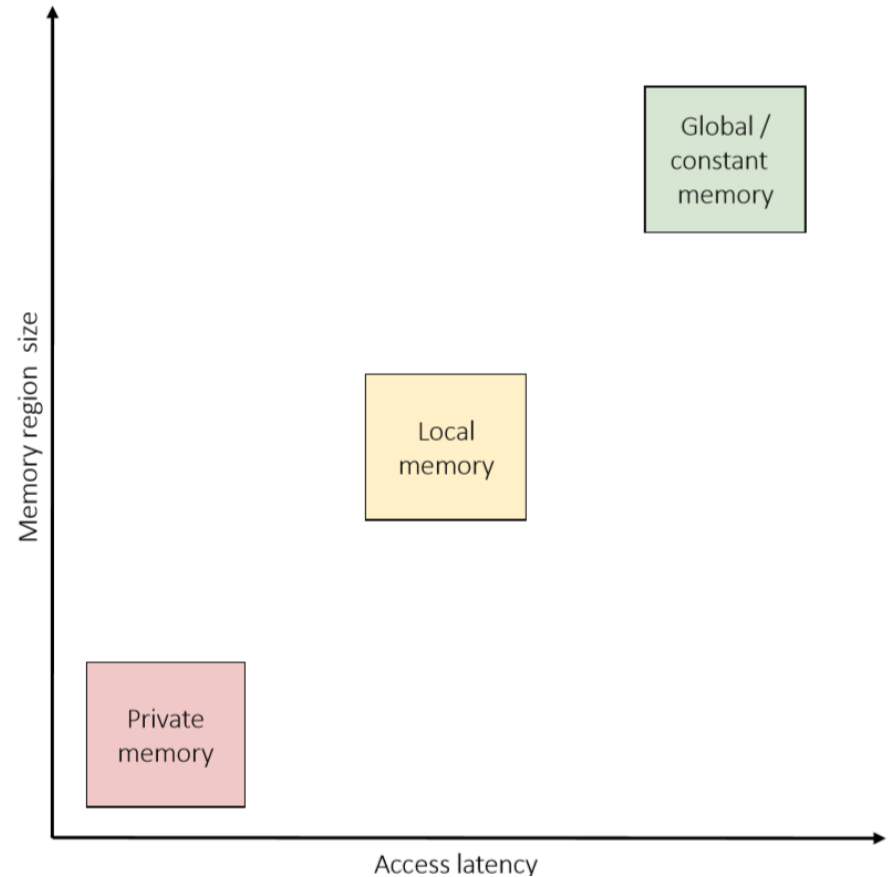
SYCL MEMORY MODEL



- Each work-item can access all of the **global memory**

SYCL MEMORY MODEL

- Each memory region has a different size and access latency
- Global memory is larger than local memory and local memory is larger than private memory
- Private memory is faster than local memory and local memory is faster than global memory



ACCESSING DATA WITH ACCESSORS

- There are a few different ways to access the data represented by an accessor
 - The subscript operator can take an **id**
 - Must be the same dimensionality of the accessor
 - For dimensions > 1, linear address is calculated in row major
- Nested subscript operators can be called for each dimension taking a **size_t**
 - E.g. a 3-dimensional accessor: `acc[x][y][z] = ...`
- A pointer to memory can be retrieved by calling **get_pointer**
 - This returns a raw pointer to the data

ACCESSING DATA WITH ACCESSORS

```
sycl::buffer<float, 1> bufA(dA.data(), sycl::range<1>(dA.size()));
sycl::buffer<float, 1> bufB(dB.data(), sycl::range<1>(dB.size()));
sycl::buffer<float, 1> bufO(dO.data(), sycl::range<1>(dO.size()));

gpuQueue.submit([&](sycl::handler &cgh) {
    sycl::accessor inA{bufA, cgh, sycl::read_only};
    sycl::accessor inB{bufB, cgh, sycl::read_only};
    sycl::accessor out{bufO, cgh, sycl::write_only};
    cgh.parallel_for<add>(sycl::range<1>(dA.size()),
        [=](sycl::id<1> i) {
            out[i] = inA[i] + inB[i];
        });
});
```

- Here we access the data of the accessor by passing in the `id` passed to the SYCL kernel function.

ACCESSING DATA WITH ACCESSORS

```
sycl::buffer<float, 1> bufA(dA.data(), sycl::range<1>(dA.size)
sycl::buffer<float, 1> bufB(dB.data(), sycl::range<1>(dB.size)
sycl::buffer<float, 1> bufO(dO.data(), sycl::range<1>(dO.size)

gpuQueue.submit([&](handler &cgh) {
    sycl::accessor inA{bufA, cgh, sycl::read_only};
    sycl::accessor inB{bufB, cgh, sycl::read_only};
    sycl::accessor out{bufO, cgh, sycl::write_only};
    cgh.parallel_for<add>(rng, [=](sycl::item<3> i) {
        auto ptrA = inA.get_pointer();
        auto ptrB = inB.get_pointer();
        auto ptrO = out.get_pointer();
        auto linearId = i.get_linear_id();

        ptrA[linearId] = ptrB[linearId] + ptrO[linearId];
    });
});
```

- Here we retrieve the underlying pointer for each of the accessors.
- We then access the pointer using the linearized id by calling the `get_linear_id` member function on the `item`.
- Again this linearization is calculated in row-major order.

QUESTIONS

EXERCISE

Code_Exercises/Data_Parallelism/source.cpp

Implement a SYCL application using `parallel_for` to add two arrays of values

- Use buffers and accessors to manage data
- Try the `sycl::range` and `sycl::nd_range` variants

