



Hands-On HPC Application Development Using C++ and SYCL

James Reinders, Phuong Nguyen, Thomas Applencourt, Rod Burns, Alastair Murray



1

HANDLING ERRORS AND DEBUGGING

LEARNING OBJECTIVES

- Learn about how SYCL handles errors
- Learn about the difference between synchronous and asynchronous exceptions
- Learn how to handle exceptions and retrieve further information
- Learn about the host device and how to use it

SYCL EXCEPTIONS

- In SYCL errors are handled by throwing exceptions.
- It is crucial that these errors are handled, otherwise your application could fail in unpredictable ways.
- In SYCL there are two kinds of error:
 - Synchronous errors (thrown in user thread).
 - Asynchronous errors (thrown by the SYCL scheduler).

HANDLING ERRORS

```
int main() {  
    queue q();  
  
    /* Synchronous code */  
  
    q.submit([&](handler &cgh) {  
  
        /* Synchronous code */  
  
        cgh.single_task<add>(buf0.get_range(), [=](id<1> i) {  
  
            /* Asynchronous code */  
  
        });  
    });  
}
```

- Kernels run asynchronously on the device, and will throw asynchronous errors.
- Everything else runs synchronously on the host, and will throw synchronous errors.

SYCL EXCEPTIONS

Synchronous
exceptions

SYCL interface

Asynchronous
exceptions

SYCL Runtime

(optional)
CPU device

Kernel
loader

Runtime
Scheduler

Data dependency
tracker

Backend interface (e.g. OpenCL API)

HANDLING ERRORS

```
class add;

int main() {
    queue q();

    /* Synchronous code */

    q.submit([&](handler &cgh) {
        /* Synchronous code */

        cgh.single_task<add>([=](id<1> i) {
            /* Asynchronous code */
        });
    }).wait();
}
```

- Code on the device runs asynchronously
- If errors are not handled, the application can fail.
- SYCL 2020 provides a default async handler that will call `std::terminate` when an asynchronous error is thrown.

```
class add;

int main() {
    std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0, 0, 0, 0 };
    try {
        queue gpuQueue(gpu_selector{});

        buffer bufA{dA};
        buffer bufB{dB};
        buffer bufO{dO};

        gpuQueue.submit([&](handler &cgh) {
            auto inA = accessor{bufA, cgh, read_only};
            auto inB = accessor{bufB, cgh, read_only};
            auto out = accessor{bufO, cgh, write_only};

            cgh.single_task<add>(bufO.get_range(), [=](id<1> i) {
                out[i] = inA[i] + inB[i];
            });
        }).wait();

    } catch (...) { /* handle errors */ }
}
```

- Synchronous errors are typically thrown by SYCL API functions.
- In order to handle all SYCL errors you must wrap everything in a try-catch block.


```

class add;

int main() {
    std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0, 0, 0, 0 };
    try{
        queue gpuQueue(gpu_selector{}, async_handler{});

        buffer bufA{dA};
        buffer bufB{dB};
        buffer bufO{dO};

        gpuQueue.submit([&](handler &cgh) {
            auto inA = accessor{bufA, cgh, read_only};
            auto inB = accessor{bufB, cgh, read_only};
            auto out = accessor{bufO, cgh, write_only};

            cgh.single_task<add>(bufO.get_range(), [=](id<1> i) {
                out[i] = inA[i] + inB[i];
            });
        }).wait();

        gpuQueue.throw_asynchronous();
    } catch (...) { /* handle errors */
    }
}

```

- Asynchronous errors that may have occurred will be thrown after a command group has been submitted to a queue.
 - To handle these errors you must provide an async handler when constructing the queue object.

to call the `throw_asynchronous` or `wait_and_throw` member functions of

```
class add;

int main() {
    std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0, 0
    try{
        queue gpuQueue(gpu_selector{}, [=](exception_list eL) {
            for (auto e : eL) { std::rethrow_exception(e); }
        });

        buffer bufA{dA};
        buffer bufB{dB};
        buffer bufO{dO};

        gpuQueue.submit([&](handler &cgh) {
            auto inA = accessor{bufA, cgh, read_only};
            auto inB = accessor{bufB, cgh, read_only};
            auto out = accessor{bufO, cgh, write_only};
```

- The async handler is a C++ lambda or function object that takes as a parameter an `exception_list`
- The `exception_list` class is a wrapper around a list of `exception_ptr`s which can be iterated over

```
int main() {  
    std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0, 0  
    try {  
        queue gpuQueue(gpu_selector{}, [=](exception_list eL) {  
            for (auto e : eL) { std::rethrow_exception(e); }  
        });  
  
        ...  
  
        gpuQueue.throw_asynchronous();  
    } catch (const std::exception& e) {  
        std::cout << "Exception caught: " << e.what()  
        << std::endl;  
    }  
}
```

- Once rethrown and caught, a SYCL exception can provide information about the error
- The what member function will return a string with more details

```
int main() {
    std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0, 0
    try {
        queue gpuQueue(gpu_selector{}, [=](exception_list eL) {
            for (auto e : eL) { std::rethrow_exception(e); }
        });

        ...

        gpuQueue.throw_asynchronous();
    } catch (const sycl::exception& e) {
        std::cout << "Exception caught: " << e.what();
        std::cout << " With OpenCL error code: "
        << e.get_cl_code() << std::endl;
    }
}
```

- In SYCL 1.2.1, if the exception has an OpenCL error code associated with it this can be retrieved by calling the `get_cl_code` member function
- If there is no OpenCL error code this will return `CL_SUCCESS`
- SYCL 2020 provides the `error_category_for` templated free function that allows for the category of the exception depending on the backend used (e.g. `backend::opencl`), and a `code()` value() will correspond to the backend error

```
int main() {
    std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0, 0, 0, 0 };

    queue gpuQueue(gpu_selector{}, [=](exception_list eL) {
        for (auto e : eL) { std::rethrow_exception(e); }
    });
    context gpuContext = gpuQueue.get_context();

    try {
        ...
        gpuQueue.wait_and_throw();
    } catch (const sycl::exception& e) {
        if (e.has_context()) {
            if (e.get_context() == gpuContext) {
                /* handle error */
            }
        }
    }
}
```

- The `has_context` member function will tell you if there is a SYCL context associated with the error
- If that returns true then the `get_context` member function will return the associated SYCL context object

EXCEPTION TYPES

- SYCL 2020 has a single `sycl::exception` type which provides different error codes
 - e.g., `errc::runtime`, `errc::kernel`

DEBUGGING SYCL KERNEL FUNCTIONS

- Top debugging tip: use CPU devices during development as much as is appropriate.
- SYCL 2020 only guarantees that a device will always be available.
- We can query the `host_debuggable` device aspect to check for host-level type debugging support. Such devices allow us to debug a SYCL kernel function using a standard C++ debugger (e.g., gdb).

```
class add;

int main() {
    std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0, 0, 0, 0 };
    try{
        queue hostQueue(aspect_selector<aspect::host_debuggable>(), async_handler{});

        buffer bufA{dA};
        buffer bufB{dB};
        buffer bufO{dO};

        hostQueue.submit([&](handler &cgh) {
            auto inA = accessor{bufA, cgh, read_only};
            auto inB = accessor{bufB, cgh, read_only};
            auto out = accessor{bufO, cgh, write_only};

            cgh.single_task<add>(bufO.get_range(), [=](id<1> i) {
                out[i] = inA[i] + inB[i];
            });
        });
        hostQueue.wait_and_throw();
    } catch (...) { /* handle errors */ }
}
```

- In general, a SYCL application can be debugged on the CPU device by switching the queue for a CPU queue
- Replacing the device selector for the `aspect_selector` will ensure that the queue submits all work to the device with the requested aspects, in this case a `host_debuggable` device

QUESTIONS

EXERCISE

Lesson_Materials/Lecture_03_Error_Handling

- Introduce a synchronous error, and
- Introduce an asynchronous error.
- Catch them and report them without aborting the program (so we get to see both error messages).
- The try/catch framework is already in place.

