

HANDLING ERRORS AND DEBUGGING

LEARNING OBJECTIVES

- Learn about how SYCL handles errors
- Learn about the difference between synchronous and asynchronous exceptions
- Learn how to handle exceptions and retrieve further information
- Learn about the host device and how to use it

SYCL EXCEPTIONS

- In SYCL errors are handled by throwing exceptions.
- It is crucial that these errors are handled, otherwise your application could fail in unpredictable ways.
- In SYCL there are two kinds of error:
 - Synchronous errors (thrown in user thread) .
 - Asynchronous errors (thrown by the SYCL scheduler).

SYCL EXCEPTIONS

Synchronous
exceptions

SYCL interface

Asynchronous
exceptions

SYCL Runtime

(optional)
CPU device

Kernel
loader

Runtime
Scheduler

Data dependency
tracker

Backend interface (e.g. OpenCL API)

HANDLING ERRORS

```
class add;

int main() {
    std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0, 0, 0, 0 };
    queue gpuQueue(gpu_selector{});
    buffer bufA{dA};
    buffer bufB{dB};
    buffer bufO{dO};

    gpuQueue.submit([&](handler &cgh) {
        auto inA = accessor{bufA, cgh, read_only};
        auto inB = accessor{bufB, cgh, read_only};
        auto out = accessor{bufO, cgh, write_only};

        cgh.single_task<add>(bufO.get_range(), [=](id<1> i) {
            out[i] = inA[i] + inB[i];
        });
    }).wait();
}
```

- If errors are not handled, the application can fail:
 - SYCL 1.2.1 application will fail silently.
 - SYCL 2020 provides a default async handler that will call `std::terminate` when an asynchronous error is thrown.

```
class add;

int main() {
    std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0, 0, 0, 0 };
    try {
        queue gpuQueue(gpu_selector{});

        buffer bufA{dA};
        buffer bufB{dB};
        buffer bufO{dO};

        gpuQueue.submit([&](handler &cgh) {
            auto inA = accessor{bufA, cgh, read_only};
            auto inB = accessor{bufB, cgh, read_only};
            auto out = accessor{bufO, cgh, write_only};

            cgh.single_task<add>(bufO.get_range(), [=](id<1> i) {
                out[i] = inA[i] + inB[i];
            });
        }).wait();
    } catch (...) { /* handle errors */ }
}
```

- Synchronous errors are typically thrown by SYCL API functions.
- In order to handle all SYCL errors you must wrap everything in a try-catch block.

```
class add;

int main() {
    std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0, 0, 0, 0 };
    try{
        queue gpuQueue(gpu_selector{}, async_handler{});

        buffer bufA{dA};
        buffer bufB{dB};
        buffer bufO{dO};

        gpuQueue.submit([&](handler &cgh) {
            auto inA = accessor{bufA, cgh, read_only};
            auto inB = accessor{bufB, cgh, read_only};
            auto out = accessor{bufO, cgh, write_only};

            cgh.single_task<add>(bufO.get_range(), [=](id<1> i) {
                out[i] = inA[i] + inB[i];
            });
        }).wait();

        gpuQueue.throw_asynchronous();
    } catch (...) { /* handle errors */
    }
}
```

- Asynchronous errors that may have occurred will be thrown after a command group has been submitted to a queue.
 - To handle these errors you must provide an `async` handler when constructing the queue object.
- Then you must also call the `throw_asynchronous` or `wait_and_throw` functions of the queue class.

```
class add;

int main() {
    std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0, 0, 0, 0 };
    try{
        queue gpuQueue(gpu_selector{}, [=](exception_list eL) {
            for (auto e : eL) { std::rethrow_exception(e); }
        });

        buffer bufA{dA};
        buffer bufB{dB};
        buffer bufO{dO};

        gpuQueue.submit([&](handler &cgh) {
            auto inA = accessor{bufA, cgh, read_only};
            auto inB = accessor{bufB, cgh, read_only};
            auto out = accessor{bufO, cgh, write_only};

            cgh.single_task<add>(bufO.get_range(), [=](id<1> i) {
                out[i] = inA[i] + inB[i];
            });
        }).wait();

        gpuQueue.throw_asynchronous();
    } catch (...) { /* handle errors */ }
}
```

- The async handler is a C++ lambda or function object that takes as a parameter an `exception_list`
- The `exception_list` class is a wrapper around a list of `exception_ptrs` which can be iterated over

`exception_ptrs` can be rethrown by passing them to


```
int main() {  
    std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0, 0, 0, 0 };  
    try {  
        queue gpuQueue(gpu_selector{}, [=](exception_list eL) {  
            for (auto e : eL) { std::rethrow_exception(e); }  
        });  
  
        ...  
  
        gpuQueue.throw_asynchronous();  
    } catch (const std::exception& e) {  
        std::cout << "Exception caught: " << e.what()  
        << std::endl;  
    }  
}
```

- Once rethrown and caught, a SYCL exception can provide information about the error
- The `what` member function will return a string with more details

```
int main() {
    std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0, 0, 0, 0 };
    try {
        queue gpuQueue(gpu_selector{}, [=](exception_list eL) {
            for (auto e : eL) { std::rethrow_exception(e); }
        });

        ...

        gpuQueue.throw_asynchronous();
    } catch (const sycl::exception& e) {
        std::cout << "Exception caught: " << e.what();
        std::cout << " With OpenCL error code: "
        << e.get_cl_code() << std::endl;
    }
}
```

- In SYCL 1.2.1, if the exception has an OpenCL error code associated with it this can be retrieved by calling the `get_cl_code` member function
- If there is no OpenCL error code this will return `CL_SUCCESS`
- SYCL 2020 provides the `error_category_for` templated free function that allows checking for the category of the exception depending on the backend used (e.g. `backend::opencl`), and `e.code().value()` will correspond to the backend error code.

```
int main() {
    std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0, 0, 0, 0 };

    queue gpuQueue(gpu_selector{}, [=](exception_list eL) {
        for (auto e : eL) { std::rethrow_exception(e); }
    });
    context gpuContext = gpuQueue.get_context();

    try {
        ...
        gpuQueue.wait_and_throw();
    } catch (const sycl::exception& e) {
        if (e.has_context()) {
            if (e.get_context() == gpuContext) {
                /* handle error */
            }
        }
    }
}
```

- The `has_context` member function will tell you if there is a SYCL context associated with the error
- If that returns true then the `get_context` member function will return the associated SYCL context object

EXCEPTION TYPES

- In SYCL 1.2.1 there are a number of different exception types that inherit from `std::exception`
 - E.g. `runtime_error`, `kernel_error`
- SYCL 2020 only has a single `sycl::exception` type which provides different error codes
 - E.g. `errc::runtime`, `errc::kernel`

DEBUGGING SYCL KERNEL FUNCTIONS

- Every SYCL 1.2.1 implementation is required to provide a host device
 - This device executes native C++ code but is guaranteed to emulate the SYCL execution and memory model
- This means you can debug a SYCL kernel function by switching to the host device and using a standard C++ debugger
 - For example gdb

- SYCL 2020 only guarantees that a device will always be available, and users can query the `host_debuggable` device aspect to check whether they can use the same functionality as the SYCL 1.2.1 host device


```

class add;

int main() {
    std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0, 0, 0, 0 };
    try{
        queue hostQueue(aspect_selector<aspect::host_debuggable>(), async_handler{});

        buffer bufA{dA};
        buffer bufB{dB};
        buffer bufO{dO};

        hostQueue.submit([&](handler &cgh) {
            auto inA = accessor{bufA, cgh, read_only};
            auto inB = accessor{bufB, cgh, read_only};
            auto out = accessor{bufO, cgh, write_only};

            cgh.single_task<add>(bufO.get_range(), [=](id<1> i) {
                out[i] = inA[i] + inB[i];
            });
        });
        hostQueue.wait_and_throw();
    } catch (...) { /* handle errors */ }
}

```

- Any SYCL application can be debugged on the host device by switching the queue for a host queue
- Replacing the device selector for the `aspect_selector` will ensure that the queue submits all work to the device with the requested aspects, in this case a host debuggable device

1.2.1, `host_selector` would be used instead, deprecated in SYCL

QUESTIONS

EXERCISE

Code_Exercises/Exercise_4_Handling_Errors/source

Add error handling to a SYCL application for both synchronous and asynchronous errors.