

# ASYNCHRONOUS EXECUTION

# LEARNING OBJECTIVES

- Learn about how commands are enqueued asynchronously
- Learn about the different reasons for synchronization
- Learn about the different ways to perform synchronization

# ASYNCHRONOUS EXECUTION

- All command submitted to a `queue` are done so asynchronously.
- The functions return immediately and the command is run in a background thread.
- This includes individual commands like `memcpy` and collections of commands derived from a command group.
- This means you have to synchronize with those commands.

# SYNCHRONIZATION

There are a number of reasons why you need to synchronize with commands

- Await completion of a kernel function.
- Await the results of a computation.
- Await error conditions which come from a failure to execute any of the commands.

# SYNCHRONIZATION WITH KERNEL FUNCTIONS

There are two ways ways to synchronize with kernel functions.

- Calling `wait` on an `event` object returned from enqueueing a kernel function command, either via a command group or a shortcut function.
- Calling `wait` or `wait_and_throw` on the `queue` itself.

# SYNCHRONIZING WITH KERNEL FUNCTIONS (BUFFERS/ACCESSORS)

```
buf = sycl::buffer(data, sycl::range{1024});

gpuQueue.submit([&](sycl::handler &cgh){
    auto acc = sycl::accessor(buf, cgh);

    cgh.parallel_for<kernel_a>(sycl::range{1024},
        [=](sycl::id<1> idx){
            acc[idx] = /* some computation */
        });
}).wait();
```

- Calling `wait` on an event object returned from enqueueing a command group will wait for the commands from that command group to complete.
- This is how we have synchronized in our examples so far.
- This effectively creates a blocking operations that will complete in place by immediately synchronizing.

# SYNCHRONIZING WITH KERNEL FUNCTIONS (BUFFERS/ACCESSORS)

```
buf = sycl::buffer(data, sycl::range{1024});

gpuQueue.submit([&](sycl::handler &cgh){
    auto acc = sycl::accessor(buf, cgh);

    cgh.parallel_for<kernel_a>(sycl::range{1024},
        [=](sycl::id<1> idx){
            acc[idx] = /* some computation */
        });
});

gpuQueue.wait();
```

- Calling `wait` or `wait_and_throw` on a queue will wait for all commands enqueued to it to complete.
- Note that command groups do not create commands to copy data back to the host application.

## SYNCHRONIZING WITH KERNEL FUNCTIONS (USM)

```
auto devicePtr = usm_wrapper<int>(
    malloc_device<int>(1024, gpuQueue));

gpuQueue.memcpy(devicePtr, data, sizeof(int)).wait();

gpuQueue.parallel_for<kernel_a>(sycl::range{1024},
    [=](sycl::id<1> idx){
        devicePtr[idx] = /* some computation */
    }).wait();
```

- Calling `wait` on an event object returned from functions such as `memcpy` or the queue shortcuts will wait for that specific command to complete.
- Again this is how we have synchronized in our examples so far.



## SYNCHRONIZING WITH KERNEL FUNCTIONS (USM)

```
auto devicePtr = usm_wrapper<int>(
    malloc_device<int>(1024, gpuQueue));

gpuQueue.memcpy(devicePtr, data, sizeof(int));

gpuQueue.wait();

gpuQueue.parallel_for<kernel_a>(sycl::range{1024},
    [=](sycl::id<1> idx){
        devicePtr[idx] = /* some computation */
    });

gpuQueue.wait();
```

- Again calling `wait` or `wait_and_throw` on a queue will wait for all commands enqueued to it to complete.
- Note you generally don't want to call `wait` on the queue after every command, instead you want to create dependencies between commands, which we cover in the next lecture.

## SYNCHRONIZING WITH DATA

There are multiple ways ways to synchronize with data, but it differs depending on the data management model you are using.

- When using the USM data management model you can synchronize the same way you would for kernel functions, calling `wait` on an `event` or the `queue`.
- When using the buffer/access data management model command groups don't automatically copy data back so there are other ways to synchronize with the data.
  - Creating a `host_accessor`.
  - Destroying the `buffer`.

## SYNCHRONIZING WITH DATA (USM)

```
gpuQueue.memcpy(data, devicePtr, sizeof(int)).wait();
```

```
gpuQueue.memcpy(data, devicePtr, sizeof(int));  
gpuQueue.wait();
```

- Simply call `wait` on the event returned from `memcpy`.
- Alternatively call `wait` on the queue.

## SYNCHRONIZING WITH DATA (BUFFER/ACCESSOR)

```
buf = sycl::buffer(data, sycl::range{1024});

gpuQueue.submit([&](sycl::handler &cgh){
    auto acc = sycl::accessor{buf, cgh};

    cgh.parallel_for<kernel_a>(sycl::range{1024},
        [=](sycl::id<1> idx){
            acc[idx] = /* some computation */
        });
});

{
    auto hostAcc = buf.get_host_access();

    hostAcc[idx] = /* some computation */
}
```

- A `host_accessor` gives immediate access to the data managed by a `buffer` in the host application.
- This will wait for any kernel functions accessing the `buffer` to complete and then copying the data back to the host.
- It will also block any other accessor accessing a `buffer` until it is destroyed.

## SYNCHRONIZING WITH DATA (BUFFER/ACCESSOR)

```
{
    buf = sycl::buffer(data, sycl::range{1024});

    gpuQueue.submit([&](sycl::handler &cgh){
        auto acc = sycl::accessor{buf, cgh};

        cgh.parallel_for<kernel_a>(sycl::range{1024},
            [=](sycl::id<1> idx){
                acc[idx] = /* some computation */
            });
    });
}
```

- A buffer will also synchronize the data it manages on destruction.
- It will wait for any kernel functions accessing it to complete and copy the data back to the origin address before completing destruction.

## SYNCHRONIZING WITH ERRORS

- Errors are handled by a `queue` and any asynchronous errors can be produced during any of the synchronization methods we've looked at.
- The best way to ensure all errors are caught is to synchronize by calling `wait` or `wait_and_throw` on the `queue`.

# QUESTIONS

## EXERCISE

`Code_Exercises/Exercise_9_Synchronization/source`

Try out the different methods of synchronizing with a kernel function and the resulting data from the computation.